

Buying retained payers, not installs

A growth diagnostic and 12-week experiment plan for Kuku TV.

Preetam Priyabrat

Hive School PGP Cohort 2

September 2026

Built on public data only

Two things changed in 2026

APRIL 2026

A free competitor arrived with distribution you cannot buy

- JioHotstar launched **Tadka**. Free, ad-supported, 100+ original shows.
- **100 million users in two months**. Watch time per viewer up fivefold.
- Launched inside an app carrying **300 million IPL viewers**.

What this changes: the fight moved from paid versus paid, to paid versus free. That raises the price of attention in the ad auction and lowers what a first-time viewer expects to pay.

JUNE 2026

The category's own operators flagged a retention problem

Getting people to try microdrama is easy. Getting them to come back is the hard part. Acquisition is not the challenge today. Retention is.

Rohan Nayak, Pocket FM, on shutting Pocket TV after five months

Worth noting: Pocket FM relaunched in the category two months later. This was a criticism of how the model is run, not an exit from the format.

THE QUESTION THIS DECK ANSWERS

Kuku grew revenue sevenfold in FY26 while moving toward break-even, so acquisition got sharply more efficient. **Can that hold now that the competition is free, and is Kuku's advertising still optimised for the event that predicts a paying customer?**

Microdrama is a paid-acquisition business with an episode library attached

80%

of short drama installs are bought, not organic. There is no organic engine to fall back on.

2.3×

the Android baseline CPI, because every microdrama app bids for the same vertical video placements.

140%

of a user's revenue can be spent acquiring them, per an Indian platform founder.

The first question I would ask inside Kuku

The category converts **2 to 5%** of viewers into payers. Kuku TV's public numbers imply roughly **13%**.

Either Kuku converts three to five times better than the market, which is a real moat worth protecting, or "paying subscriber" is counting **active mandates** rather than **people getting value**.

I do not know which. The answer changes the entire growth plan.

What this means for the job

When four fifths of installs are bought, CPIs carry a category premium, and the plan price sits close to acquisition cost, the biggest lever is not how much you spend.

It is which user the spend selects for.

That is a measurement and event-optimisation problem before it is a budget problem.

Kuku has already run this play once

A published case study, with a named Kuku growth marketer and a rupee figure attached.

₹26 lakh

Marketing spend recovered after Kuku FM blocked fake installs, click-spamming affiliates and reward-chasing traffic.

Junk installs were draining our resources without delivering returns. We needed to filter our ad traffic and optimise campaigns better.

Aashita Jain, Growth Marketing Manager, Kuku FM

- **The symptom:** install charts climbing, post-install activity flat.
- **The causes:** junk APK installs, affiliates faking clicks to steal attribution, and users installing only to claim a reward.
- **The fix:** clean the traffic before it reaches attribution. Fees dropped, and channel performance became readable for the first time.

What this deck proposes next

Kuku proved **an install is not a user**, and fixed it at the traffic layer.

The same logic applies one step deeper: **a signup is not a retained payer**. That fix sits at the optimisation-event layer.

Why I open here rather than with a critique. Telling a team its growth model is broken rests on data I do not have. Saying you already proved this principle once, here is where it applies next is the same analysis, it is verifiable, and it starts from something the team got right.

What I think is happening, stated as a hypothesis

Kuku's paid ads may be buying users who convert quickly but do not survive to their first renewal. Optimise the ads against an event that actually predicts a 30-day payer, and **cost per retained payer falls without losing install volume**.

I cannot see Kuku's current optimisation event from outside. Confirming it is Week 1, Day 1, not an assumption this deck rests on.

SIX THINGS TO TEST, IN ORDER

	HYPOTHESIS	LAYER	RESOLVED BY
H1	Optimising deeper in the funnel lowers cost per 30-day retained payer	Acquisition	Experiment 1
H2	Episodes watched before the paywall predicts retention better than any targeting variable	Activation	Cohort analysis, Week 2
H3	Value varies more by which show a user enters on than by who was targeted	Content	Experiment 2
H4	Cost and value flip across languages and city tiers, so one blended target misallocates budget	Regional	Dashboard, Week 3
H5	A small paid step before the subscription ask increases retained payers	Monetisation	Experiment 3
H6	Much of churn is failed payments, recoverable far cheaper than buying a new user	Recovery	Week 1, no-regret

The growth system, and the metric that governs each step

Everything in this deck attaches to one of these fourteen steps. Nothing sits outside it.

ACQUISITION

01
Ad
impression

02
Click

03
Install

04
Registration

CPM

CTR
CPC

CPI

Reg.
rate

ACTIVATION

05
First story
started

06
Episode 3

07
Episode 10

08
Paywall
exposed

Activation
rate

Ep-3
depth

Ep-10
depth

Paywall
exposure

MONETISATION & RETENTION

09
Trial
started

10
First
payment

11
D7 retained

12
D30
retained

13
Renewal

14
LTV

Trial
CVR

Paid
CVR

D7

D30

Renewal
rate

LTV
ROAS

The metric that governs everything

Cost per 30-day retained payer

CPI and trial CAC are inputs. This is the number the whole system steers by.

Why the funnel is drawn this wide

Steps 05 to 08 are where habit forms, and the only steps content decisions can move. Most decks collapse them and lose the content lever entirely.

The catch on going deeper

A deeper event predicts value better but fires less often. Too deep and the algorithm never exits learning. Signal quality against signal density.

Three places value is most likely leaking

Each maps to a stretch of the funnel, and each has one measurement that proves or kills it.

1 · Selection

STEPS 01 TO 04

Ads optimised for the earliest conversion find the cheapest people who will trigger it, not the people who renew.

Measure: report every campaign twice, once on trial CAC and once on cost per 30-day payer. If the rankings differ, the leak is real.

2 · Habit

STEPS 05 TO 08

If the paywall lands before enough episodes, the user is asked to commit before the product has proved itself. With a free alternative now at 100 million users, that ask is harder than it was last year.

Measure: retention by episodes watched before the paywall. A steep curve means timing beats price.

3 · Collection

STEPS 11 TO 14

UPI autopay succeeds about 85% in month one and drops to around 70% by month six. Payers who wanted to stay get counted as churn.

Measure: split churn into people who cancelled and people whose payment failed. These users are already paid for.


The rule I would hold the whole programme to. Growth must come from people choosing to stay, not from making cancellation hard. That is a commercial position before it is an ethical one: revenue built on friction does not survive a free competitor, and it turns into refunds, bad reviews and cancelled mandates that push acquisition cost back up.

How I would measure it

No optimisation event changes until we can prove it predicts a paying customer 30 days later.



THE EVENTS WE TRACK



■ Candidate optimisation events ■ Usually missing

The two events most teams never connect

renewal_failure and **autopay_revoked** usually live only in the billing system. Because they never reach the marketing data, failed payments stay invisible to the growth team, and CAC gets calculated against a payer base the business is about to lose for reasons unrelated to intent.

Piping these two events into the warehouse is the cheapest, highest-value change on this deck.

THREE RULES I WOULD RUN THE SYSTEM BY

An event becomes a bidding target only if it **predicts value** and **fires often enough** to exit learning.

Every campaign carries **two CAC numbers** side by side. Any that ranks differently between them gets reviewed.

Fraud filtering sits **before** the MMP, not after. Kuku has run this play before.

What I would actually run in Ads Manager

CAMPAIGN STRUCTURE

Broad prospecting
Lookalike built on 30-day payers , not installers
Lookalike on top-spending payers
One campaign per language
Campaigns for the highest-value shows
Retargeting: installed but never watched

Why languages run separately

Blend them and you blend their learning. If cost and value move in opposite directions across languages, a blended target quietly overfunds the cheapest install.

CREATIVE

HOOK	TESTS FOR
Cliffhanger	Curiosity
Emotional payoff	Genre fit
Character-led	Series loyalty
Creator-style UGC	Feed-native intent
Offer-forward	Price sensitivity

The creative test I would run early

Offer-led and story-led hooks can produce the same CPI and very different retention. The creative that wins the click is not always the one that wins the renewal.

WHICH EVENT WE OPTIMISE FOR

Install	Most signal, least meaning
Registration	Still before any value
First episode watched	First real intent
Episode 3 or 10	My starting bet
First payment	Strong link, lower volume
30-day payer	Best link, too rare to bid on

Pick the **deepest event that still fires often enough daily**. Use cost per 30-day payer to read results, not to bid.

Three experiments, run in sequence

Never at the same time on the same audience, or no result is attributable.

Move the optimisation event

WEEKS 5–6 · H1

TEST

Same creative, same audience, same budget.
Only the optimisation event changes.

PRIMARY METRIC

Cost per 30-day retained payer.

KILL IF

Worse than control at readout, or install volume drops below the floor for seven straight days.

Move the paywall later

WEEKS 7–8 · H2, H3

TEST

Paywall triggers after more episodes, held constant across shows so the content effect stays separable.

PRIMARY METRIC

30-day retained payers per 1,000 installs.

KILL IF

Revenue per install falls without a matching retention gain at day 30.

Add a small paid step

WEEKS 9–10 · H5

TEST

A small optional unlock before the subscription ask. The rail exists already: Kuku Coins are live, and pay-per-piece revenue grew from ₹0.14 Cr to ₹12 Cr in FY25.

PRIMARY METRIC

30-day retained payers per 1,000 installs.

KILL IF

It eats subscription conversion, or revenue concentrates in a thin top cohort.

Running alongside all three, from week one

Connect **renewal_failure** and **autopay_revoked** to the warehouse, split churn into cancelled versus failed, and move payment reminders to WhatsApp. Not an experiment, because winning back customers who wanted to stay is never the wrong move.

If I joined on Monday

Diagnosis before experiments. Nothing launches until the measurement can read it.

DAYS 1 TO 3

Get the ground truth

- Access to Meta, Google, YouTube, the MMP and the warehouse
- Pull 90 days of spend by channel, campaign and creative
- Find out **what event each channel is currently optimised for**
- Confirm what "paying subscriber" actually counts

DAYS 4 TO 7

Build the view that does not exist

- Join spend to revenue in one table
- Show trial CAC and cost per 30-day payer side by side
- List every campaign whose ranking changes between the two
- Split churn into cancelled versus failed payment

WEEK 2

Name the constraint, then commit

- Pick the biggest constraint from the data, not from this deck
- Ship the failed-payment recovery work
- Design Experiment 1 and set the reporting cadence
- Agree kill criteria with data and finance **before** launch

WEEKS 3 TO 12

WEEKS 3–4

Segment economics established. Experiment 1 built.

WEEKS 5–6

Experiment 1 live. No interim decisions.

WEEKS 7–8

Experiment 2 live. Cohort 1 maturing.

WEEKS 9–10

Experiment 3 live. Result 1 available.

WEEK 11

Full readout across all three.

WEEK 12

Scale or kill, against week-2 criteria.

What I know, and what I would need on day one

VERIFIED FROM PUBLIC SOURCES

FY24 and FY25 financials from RoC filings · FY25 marketing spend of ₹285 Cr against ₹242 Cr revenue · FY26 revenue near ₹1,400 Cr · Kuku TV pricing and plan structure · the failed-renewal step-down mechanism in Kuku's own terms · Kuku Coins and pay-per-piece growth · the ₹26 lakh install-fraud recovery · Tadka's launch and scale · category CPI and conversion benchmarks.

SOURCES

Financials and IPO. Entrackr and Inc42 on Kuku FM FY24 and FY25 RoC filings. Startuppedia on FY25 revenue composition. Whalesbook, IndiaIPO, Exchange4media and Outlook Business on the June 2026 confidential filing, FY26 revenue and use of proceeds.

Kuku product mechanics. Kuku FM terms, conditions and support pages covering plans, refunds, Kuku Coins and the Auto-Pay Not Realised renewal step-down.

Prior acquisition remediation. mFilterIt case study on Kuku FM, including the ₹26,11,280 recovered-spend figure and the quotation from Aashita Jain.

UNKNOWN WITHOUT INTERNAL ACCESS

The current optimisation event per channel · CAC by channel and creative · FY26 marketing spend · the definition behind "paying subscriber" · D7 and D30 retention · LTV by cohort and language · cancelled versus failed-payment churn · episodes watched before the paywall · app store versus web checkout mix.

No Kuku internal figure has been estimated anywhere in this deck.

Competition. Variety, The Hollywood Reporter and Indian Television on Tadka. Storyboard18 and Business Standard on the wider microdrama field. Storyboard18 and BestMediaInfo on Pocket FM.

Category economics. Business of Apps on short drama CPI indices. Mistplay on paid-acquisition dependence. Flicknexs on monetisation mechanics. BestMediaInfo on Indian microdrama acquisition costs.

Payments and retention. Razorpay and PayU on UPI AutoPay mechanics, renewal decay and involuntary churn. Business Standard on national mandate revocation volumes.